TRADEPULSE

Report B — Critical & Logical Judge Q&A; Fire Drill

Guidebook: every panel asks some version of three questions. Prepare real answers — not deck slides.

0. The three questions you should expect

Q0.1 — Who is your first paying customer, specifically, and how do you get to them?

Answer frame: First buyer = Head of Trade Finance Operations at a GIFT City IBU (or a mid-size private bank trade desk with IFSC corridor volume). They feel exception backlog and audit pressure. Path: warm intro via IFSC/trade-ops community, 4–6 week pilot on one documentary profile (invoice + BoL), success metric = fewer unexplained exceptions + complete examiner packs — not "AI approval rate." Champion: ops head; blocker: compliance/model-risk review — we position as decision support with HITL.

Q0.2 — What part of this would break first if you tried to sell it to a real bank/NBFC next week?

Answer frame: (1) Enterprise IAM/SSO, VPC private networking, and change-control — not the happy-path demo. (2) Model-risk / vendor due diligence on Bedrock prompts and audit retention. (3) Live sanctions/list providers and data contracts beyond fixtures. (4) Latency/SLA of synchronous Textract+Bedrock on large PDFs. We would sell a supervised pilot with fixture-labeled adapters and a human-in-the-loop policy — not a production autopilot.

Q0.3 — If you had to start this company Monday, what would you build first — and is it what you built tonight?

Answer frame: Monday #1 = the examiner case pack + identity ladder + failure states that never false-PASS — the trust core. Tonight's hackathon build proves the agentic pipeline and live AWS path as a **prototype / proof point**. We would defer polish marketing and multi-tenant billing; we would accelerate SSO, async jobs, and one real bank policy pack. Good founders know prototype ≠ product — tonight is the proof point.

1. Problem ownership & customer specificity

Q — Why you — not a generic OCR vendor?

We own the failure taxonomy banks care about: unavailable data, candidate vs confirmed identity, bounded agents, maker–checker. OCR is commodity; examiner workflow + policy honesty is not.

Q — Who pays vs who uses?

User: documentary examiner / trade ops officer. Buyer: Head of Trade Ops / Transaction Banking. Influencer: Compliance / model risk. We don't sell to "fintechs" as a blob.

Q — Isn't this just another LC checking tool?

LC checkers chase clause matching. We chase cross-document evidence, identity ladder, and explicit non-PASS states — closer to TBML-aware ops support than clause bingo.

Q — Why GIFT / IFSC specifically?

Concentrated cross-border documentary volume + regulated posture + buyer title you can name. Domestic GSTIN-heavy flows are adjacent but not our first wedge.

2. Why now (timing)

Q — Why isn't "fintech is growing" enough?

Judges want concrete timing: agentic LLM cost curves, LEI/vLEI maturity, IFSC trade-finance posture — not GDP slides.

Q — What if rates fall / trade slows?

Exception cost per case remains; banks cut headcount before they cut audit standards. Decision support ROI rises when volume is flat but scrutiny is high.

3. Technical & architecture fire questions

Q — Where does the LLM get to decide compliance?

Nowhere. LLM proposes/challenges fields. Deterministic policy + human review decide. Consensus is confidence, not legality.

Q — What if agents disagree after 3 rounds?

REVIEW_REQUIRED. We never average conflicting values or hide disagreement.

Q — Fuzzy name match — do you block the customer?

No. Potential entity match → review. LEI-compatible evidence can support identity; vLEI is separate evidence. No auto-fraud claim.

Q — Sanctions hit on screen — are they sanctioned?

Candidate only until authoritative list evidence + configured policy. We do not announce "sanctioned" from fuzzy text.

Q — What AWS services and why Fargate?

ECS Fargate + ALB + ECR for demo speed; S3 docs; Textract OCR; Bedrock agents; IAM task roles. Amplify deferred (GitHub OAuth). Scale later via shared ALB, async workers, tighter networking.

Q — What's the monthly infra cost?

Always-on demo envelope ~USD 40–60/month mainly ALB+Fargate; Bedrock/Textract variable with usage. Tear down stacks when idle.

Q — How do you prevent DATA_UNAVAILABLE becoming PASS?

Typed outcomes in contracts; policy layer rejects missing evidence; UI shows unavailable states; QA rejects builds that soft-pass gaps.

Q — Is vLEI real in the demo?

Fixture path is labeled SYNTHETIC_DEMO_CREDENTIAL / VERIFIED_FIXTURE. Live verification only via trusted adapter when configured. Plain LEI string ≠ vLEI.

Q — Do you track the physical container?

No. Documentary reconciliation only. We never claim physical goods verification or Customs LEO filing.

Q — Single point of failure in the demo?

Synchronous Bedrock/Textract latency; ALB HTTP (no custom TLS domain yet); reliance on public ALB. Monday fix: async jobs + HTTPS domain + staging account.

4. Market, competition, business logic

Q — Who competes?

Bank in-house tools, OCR/RPA vendors, trade-doc AI startups, big-tech document AI. We differentiate on identity ladder + non-false-PASS discipline + examiner pack, not on raw extraction BLEU scores.

Q — What's the wedge pricing?

Pilot: fixed fee for one desk / corridor. Then per-case processing + seats. Enterprise: VPC deploy + support. Avoid usage-only pricing that incentivizes silent auto-pass.

Q — Regulatory risk of your product claims?

High if we claim clearance. We market decision support; preserve audit; human remains accountable. That is a feature for bank buyers.

Q — Can an NBFC buy this?

Yes if they run documentary trade/factoring ops — but first ICP remains bank/IBU trade ops with maker–checker culture.

Q — What's your unfair advantage in 12 months?

Corridor-specific rule packs + audit replay library + identity evidence quality — compounding data/policy assets, not a single prompt.

5. Coachability & Round 2 traps

Q — What did Round 1 tell you that you changed?

Prepare a real example: narrowed buyer to GIFT IBU Trade Ops head; added identity ladder & examiner pack; deployed live AWS instead of laptop-only. Do not only add a new title slide.

Q — Why should we believe you'll listen to customers?

Cite ownership rules and contract stop-ship behavior — we already refuse to invent product truth when PRD and design conflict.

Q — If we funded you, what would you stop building?

Stop: marketing surface area, extra document types without a paying corridor, unlabeled synthetic "live" claims. Start: SSO, async pipeline, one production policy pack.

6. Hostile / gotcha questions (stay calm)

Q — So the AI can approve a trade?

No. Full stop. Decision support only.

Q — Show me you didn't fabricate a sanctions source.

Walk provenance fields / adapter labels; if fixture, say fixture. Never invent a registry response.

Q — Your demo hung on create case — is it vapor?

Own it: cold Bedrock/Textract path; mitigation = pre-warmed case + status UX. Reliability is on the Monday list.

Q — Isn't agentic multi-agent just buzzword bingo?

Roles are bounded and testable; debate capped; arbiter can't invent evidence. If that's buzz, the tests would be empty — they aren't.

Q — Why shouldn't the bank build this in-house?

They can. We sell time-to-audit-grade workflow + maintained adapters. Banks buy focus, not the impossibility of building software.

Q — What happens when GLEIF is down?

IDENTITY_SOURCE_UNAVAILABLE — never auto-verify. Examiner sees source gap.

Q — Are you in the Hybrid / Quantum track?

No. Track 1 Agentic AI. We don’t claim quantum or hybrid compute.

7. 60-second closing if Q&A; runs long

TradePulse is examiner decision support for documentary trade — live on AWS, agentic but bounded, identity-honest, and allergic to false PASS. First customer: Trade Ops head at a GIFT IBU. Monday build: trust core + bank pilot plumbing. Tonight is the proof point, not the finished company.

End of Report B. Rehearse answers out loud; do not read slides.